

DynamoDB Hands-On Lab Runbook

Introduction

Goals: This runbook guides you through the process of setting up and managing Amazon DynamoDB using AWS CloudShell. The goals of this simulation are:

1. To create DynamoDB tables.
2. To load sample data into the tables.
3. To enable Point-In-Time Recovery (PITR) backups.
4. To simulate on-demand backups.

What You Will Accomplish

By the end of this simulation, you will:

- Have DynamoDB tables populated with sample data.
- Understand how to enable and manage backups for DynamoDB tables.
- Use AWS CloudShell for running AWS CLI commands to interact with DynamoDB.

Using AWS CloudShell

AWS CloudShell is a browser-based shell environment that allows you to run AWS CLI commands directly from the AWS Management Console. Unlike AWS Cloud9, which is an integrated development environment, CloudShell is ideal for quick command-line tasks without needing to set up a development environment. Using the CLI for steps 1-6 provides automation, repeatability, and ease of execution, especially when dealing with multiple resources.

Scenario: E-Commerce Platform

You are developing an e-commerce platform where users can browse a catalog of products, participate in forums, discuss specific threads, and post replies. Each type of interaction requires structured data storage, and DynamoDB offers a scalable solution for handling these requirements.

Table Breakdown and Explanation:

1. **ProductCatalog table:** Stores details about the products available for purchase. Each product has a unique ID and attributes such as name, price, and category. Used to display products on the e-commerce platform. Allows customers to view product details and make purchases.
2. **Forum table:** Manages the forums where users can create and participate in discussions. Each forum has a unique name.
3. **Thread table:** Represents individual discussion topics within forums.

4. **Reply table:** Stores replies to threads. Each reply has a unique ID and a timestamp to indicate when it was posted.

Steps to Follow

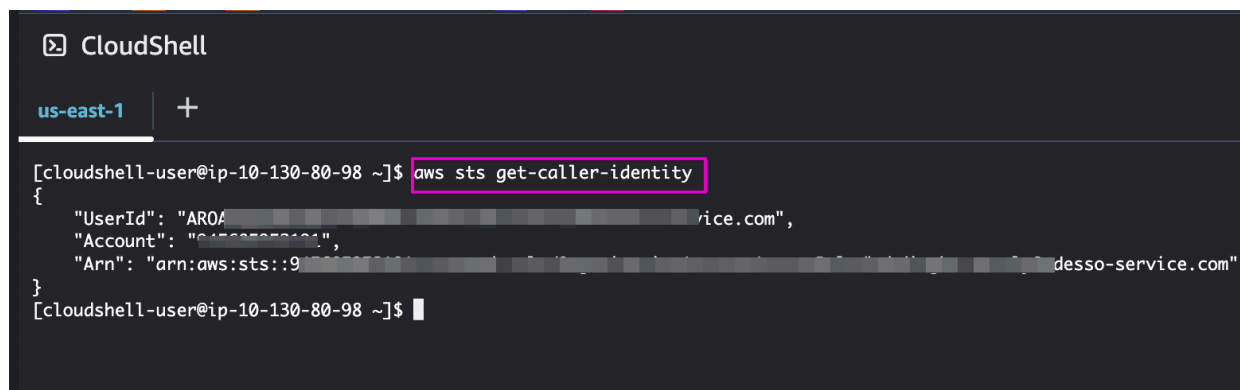
Step 1: Open AWS CloudShell

- o In the AWS Management Console, search and click on the **CloudShell** icon
- o CloudShell will open in a new browser tab.

Step 2: Verify AWS Credentials

- o In the CloudShell terminal, run the following command to verify your credentials:

\$ aws sts get-caller-identity



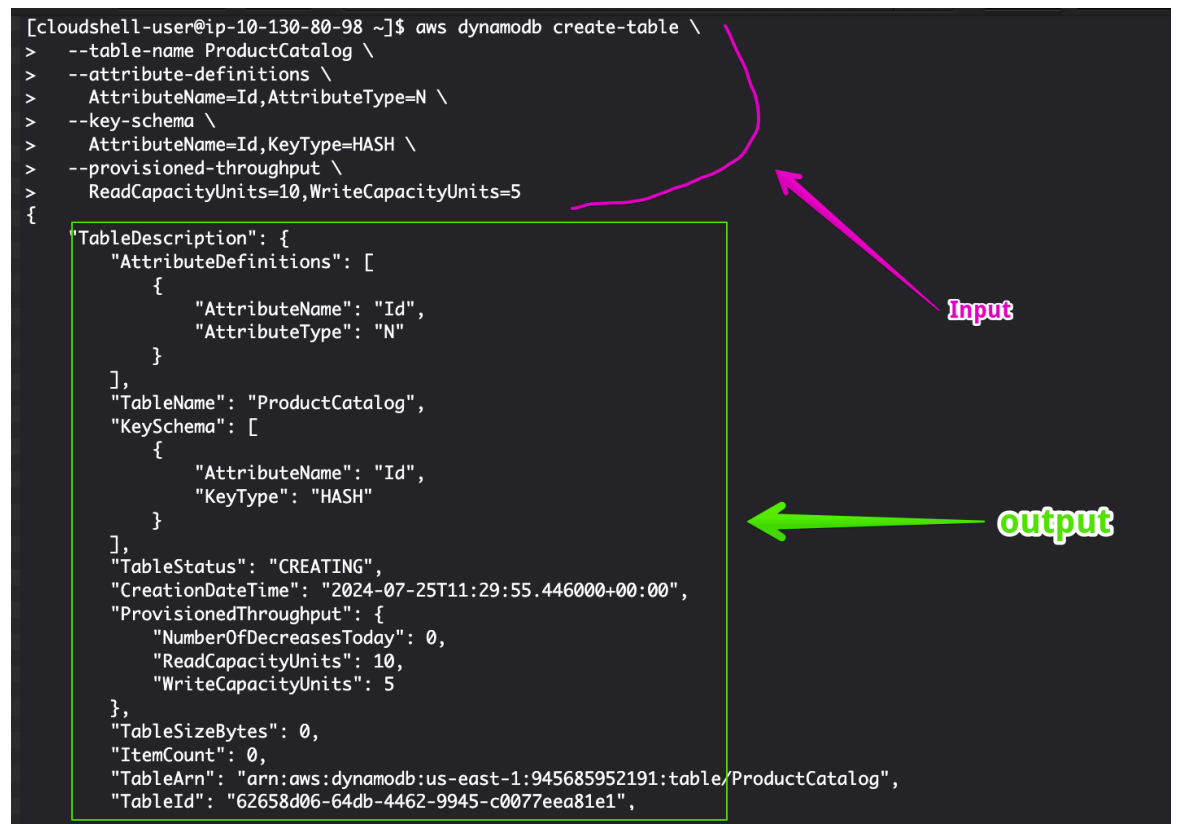
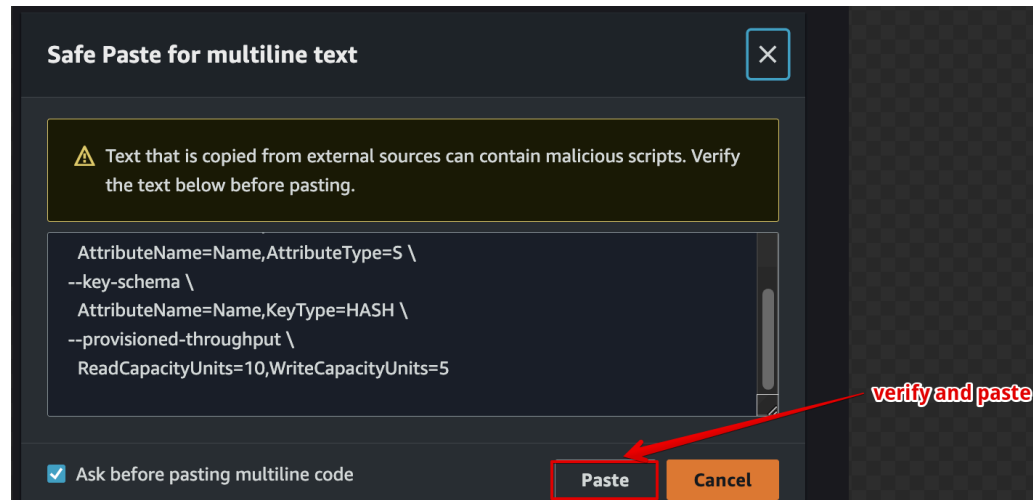
```
CloudShell
us-east-1 +
[cloudshell-user@ip-10-130-80-98 ~]$ aws sts get-caller-identity
{
  "UserId": "AROIA[REDACTED]rice.com",
  "Account": "615555555555",
  "Arn": "arn:aws:sts:9[REDACTED]desso-service.com"
}
[cloudshell-user@ip-10-130-80-98 ~]$
```

- o This command outputs details of the IAM user or role that you are using in CloudShell as shown above. Ensure the user or role has the necessary permissions to create and manage DynamoDB tables and perform backup operations.

Step 3: Create DynamoDB Tables

- o In CloudShell, create the ProductCatalog table using the AWS CLI:

```
aws dynamodb create-table \
  --table-name ProductCatalog \
  --attribute-definitions \
    AttributeName=Id,AttributeType=N \
  --key-schema \
    AttributeName=Id,KeyType=HASH \
  --provisioned-throughput \
    ReadCapacityUnits=10,WriteCapacityUnits=5
```



2. Create the Forum Table:

- o Create the Forum table using the AWS CLI:

```
aws dynamodb create-table \  
--table-name Forum \  
--attribute-definitions \  
  AttributeName=Name,AttributeType=S \  
--key-schema \  
  AttributeName=Name,KeyType=HASH \  
--provisioned-throughput \  
  ReadCapacityUnits=10,WriteCapacityUnits=5
```

3. Create the Thread Table:

- o Create the Thread table using the AWS CLI:

```
aws dynamodb create-table \  
--table-name Thread \  
--attribute-definitions \  
  AttributeName=ForumName,AttributeType=S \  
  AttributeName=Subject,AttributeType=S \  
--key-schema \  
  AttributeName=ForumName,KeyType=HASH \  
  AttributeName=Subject,KeyType=RANGE \  
--provisioned-throughput \  
  ReadCapacityUnits=10,WriteCapacityUnits=5
```

4. Create the Reply Table:

- o Create the Reply table using the AWS CLI:

```
aws dynamodb create-table \  
--table-name Reply \  
--attribute-definitions \  
  AttributeName=Id,AttributeType=S \  
  AttributeName=ReplyDateTime,AttributeType=S \  
--key-schema \  
  AttributeName=Id,KeyType=HASH \  
  AttributeName=ReplyDateTime,KeyType=RANGE \  
--provisioned-throughput \  
  ReadCapacityUnits=10,WriteCapacityUnits=5
```

5. Wait for Tables to be Created:

- o Ensure all tables are created before proceeding:

```
aws dynamodb wait table-exists --table-name ProductCatalog && \  
aws dynamodb wait table-exists --table-name Reply && \  
aws dynamodb wait table-exists --table-name Forum && \  
aws dynamodb wait table-exists --table-name Thread
```

6. Verify Table Creation:


- Confirm that all 4 tables have been created but are currently empty with no items:

```
aws dynamodb describe-table --table-name ProductCatalog
```

```
aws dynamodb describe-table --table-name Forum
```

```
aws dynamodb describe-table --table-name Thread
```

```
aws dynamodb describe-table --table-name Reply
```

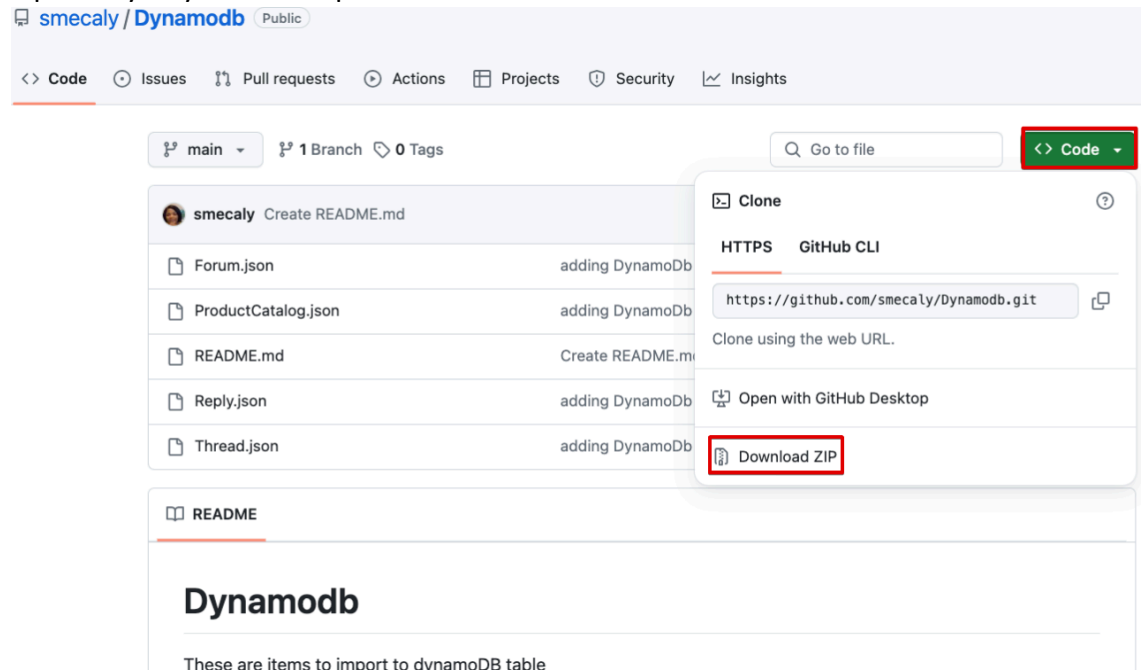


DynamoDB > Tables									
Tables (4) Info									
<input type="text" value="Find tables by table name"/> <input type="text" value="Any tag key"/> <input type="text" value="Any tag value"/> < 1 >									
<input type="checkbox"/>	Name	Status	Partition key	Sort key	Indexes	Deletion protection	Read capacity mode	Write capacity mo...	
<input type="checkbox"/>	Forum	Active	Name (S)	-	0	Off	Provisioned (10)	Provisioned (5)	
<input type="checkbox"/>	ProductCatalog	Active	Id (N)	-	0	Off	Provisioned (10)	Provisioned (5)	
<input type="checkbox"/>	Reply	Active	Id (S)	ReplyDateTime (S)	0	Off	Provisioned (10)	Provisioned (5)	
<input type="checkbox"/>	Thread	Active	ForumName (S)	Subject (S)	0	Off	Provisioned (10)	Provisioned (5)	

Step 4: Download Sample Data

1. Download Sample Data:

- Go to <https://github.com/smecaly/Dynamodb> and download the ZIP file from the repository to your desktop.



smecaly / Dynamodb Public

<> Code Issues Pull requests Actions Projects Security Insights

main 1 Branch 0 Tags

Go to file

Clone

HTTPS GitHub CLI

<https://github.com/smecaly/Dynamodb.git>

Clone using the web URL.

Open with GitHub Desktop

Download ZIP

smecaly Create README.md

- Forum.json adding DynamoDb
- ProductCatalog.json adding DynamoDb
- README.md Create README.m
- Reply.json adding DynamoDb
- Thread.json adding DynamoDb

README

Dynamodb


These are items to import to dynamoDB table

- Extract the ZIP file to access the sample data files.

Step 5: Upload Sample Data to CloudShell

1. Upload Data to CloudShell:

- o In the AWS Management Console, open CloudShell.
- o Use the upload functionality in CloudShell to upload the sample data files from your local desktop.



- o Example: Click on the upload button and select the data folder and the extracted files.

```
[cloudshell-user@ip-10-130-80-98 ~]$ ls -lh
total 24K
-rw-r--r--. 1 cloudshell-user cloudshell-user 586 Jul 25 11:59 Forum.json
-rw-r--r--. 1 cloudshell-user cloudshell-user 8.9K Jul 25 11:59 ProductCatalog.json
-rw-r--r--. 1 cloudshell-user cloudshell-user 2.2K Jul 25 12:00 Reply.json
-rw-r--r--. 1 cloudshell-user cloudshell-user 3.9K Jul 25 12:00 Thread.json
[cloudshell-user@ip-10-130-80-98 ~]$
```

Step 6: Load Sample Data into DynamoDB Tables

1. Load Data into the ProductCatalog Table:

- o Use the AWS CLI to load data into the ProductCatalog table:

```
aws dynamodb batch-write-item --request-items file://ProductCatalog.json
```

```
[cloudshell-user@ip-10-130-80-98 ~]$ aws dynamodb batch-write-item --request-items file://ProductCatalog.json
{
  "UnprocessedItems": {}
}
[cloudshell-user@ip-10-130-80-98 ~]$
```

2. Load Data into the Forum Table:

- o Use the AWS CLI to load data into the Forum table:

```
aws dynamodb batch-write-item --request-items file://Forum.json
```

3. Load Data into the Thread Table:

- o Use the AWS CLI to load data into the Thread table:

```
aws dynamodb batch-write-item --request-items file://Thread.json
```


4. Load Data into the Reply Table:

- o Use the AWS CLI to load data into the Reply table:

```
aws dynamodb batch-write-item --request-items file://Reply.json
```

5. Verify Data Load:

- o After each data load, ensure there are no unprocessed items. You should see:



Step 7: Enable Point-In-Time Recovery (PITR) Backups

What is Point-In-Time Recovery (PITR)?

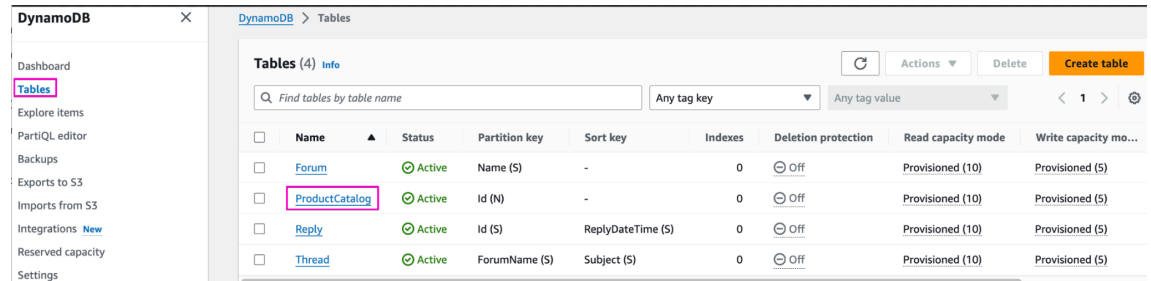
Point-In-Time Recovery (PITR) provides continuous backups of your DynamoDB table data. It allows you to restore your table to any second in time within the retention period, up to 35 days.

The goal is to ensure that your DynamoDB tables can be restored to any point in time within the last 35 days.

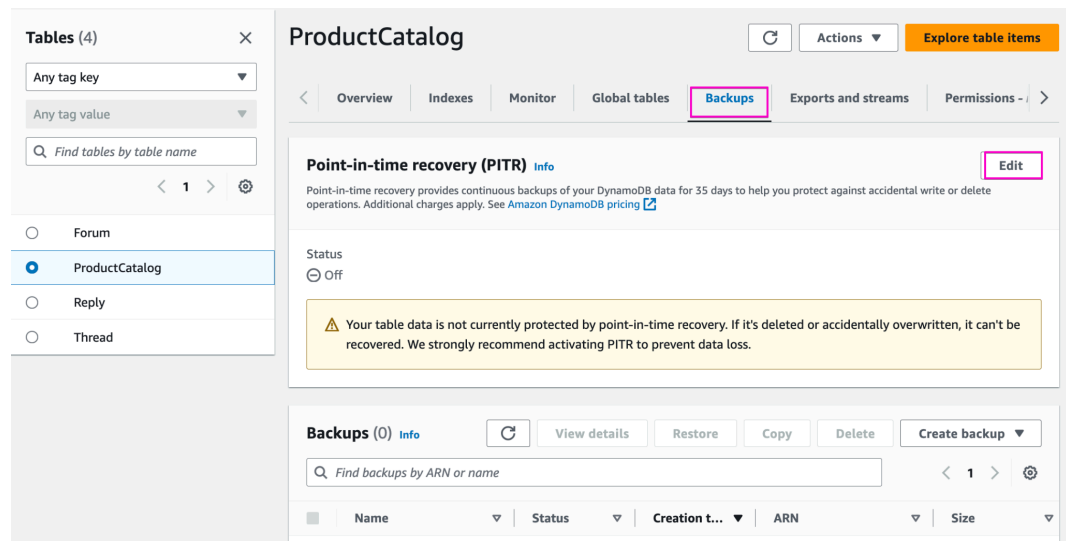
Steps to Enable PITR

1. Enable PITR via the Console:

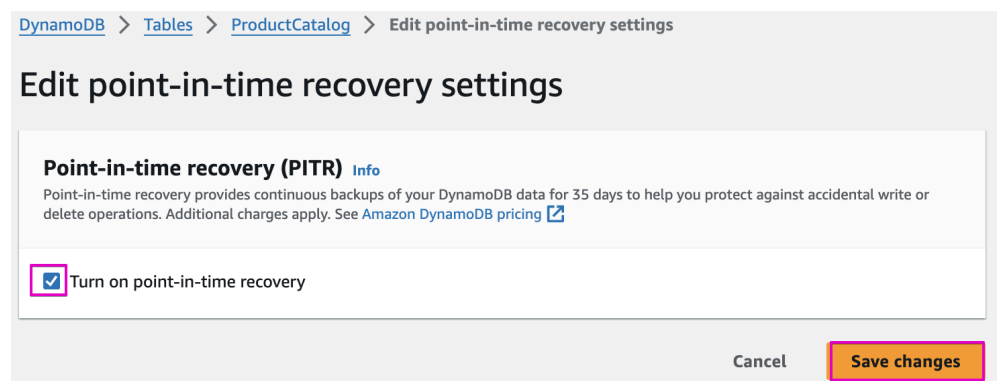
- o Go to the [DynamoDB Console](#).
- o Select the **ProductCatalog** table.



- o Navigate to the **Backups** tab and click on **Edit**.



- o In the **Point-in-time recovery settings** section, check the box to *Turn on point-in-time recovery*, and **save changes**



- o Repeat these steps for the **Forum**, **Thread**, and **Reply** tables.

Step 8: Simulate On-Demand Backups

What are On-Demand Backups?

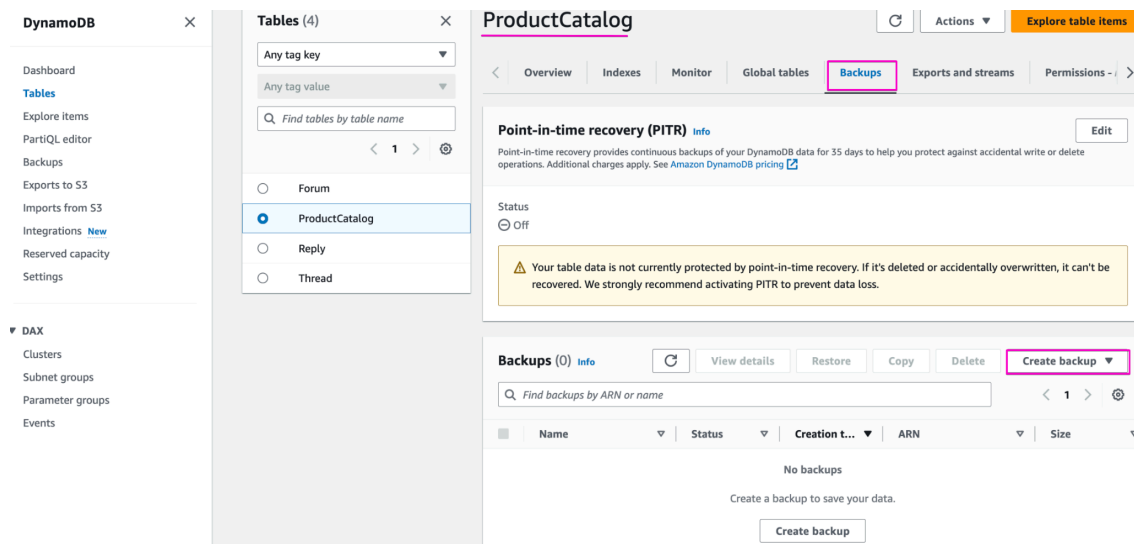
On-demand backups allow you to create full backups of your DynamoDB table data for long-term retention and archival purposes.

The goal is to create and verify on-demand backups of your DynamoDB tables.

Steps to Create an On-Demand Backup

1. Create On-Demand Backup via the Console:

- o Go to the [DynamoDB Console](#).
- o Select the **ProductCatalog** table.
- o Navigate to the **Backups** tab.
- o Click on **Create backup**.



- o select **Create on-demand backup**
- o Leave default settings to use AWS defaults or customize settings as per requirements.
- o Add a tag to the backup (e.g., **Name: ProductCatalogBackup**).
- o Click **Create**.

Backup settings [Info](#)

A backup name will be created automatically.



Default settings

Create a backup that stays in warm storage.



Customize settings

Create a backup that can transition to cold storage and be deleted as it ages.

Backup window

Start in 1 hour

Backup management

AWS Backup

Transition to cold storage

Never

Retention period

Always

Backup vault

Default

IAM Role

AWSBackupDefaultServiceRole

Tags - *optional*

AWS Backup copies tags from the DynamoDB table to the recovery point upon creation. You can specify additional tags to add to the recovery point.

Key

Value - *optional*



Name



ProductCatalogBackUp



Remove

Add new tag

You can add 49 more tags.

Cancel

Create backup

- o Repeat these steps for the **Forum**, **Thread**, and **Reply** tables (e.g., ForumBackup, ThreadBackup, ReplyBackup).
- 2. **Verify the Backup:**
 - o After the BackUp window, Ensure the backups appear in the list of backups with the status **Available**.